

Analysis of Yelp Business Intelligence Data

Part 1

Installation and Initial Setup

```
In [1]: from pyspark.sql import SparkSession
spark = SparkSession \
    .builder \
    .appName("Python Spark SQL basic example") \
    .config("spark.some.config.option", "some-value") \
    .getOrCreate()
sc.install_pypi_package("pandas==1.0.3")
sc.install_pypi_package("matplotlib==3.2.1")
sc.install_pypi_package("seaborn==0.10.0")
```

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
2	application_1606117717112_0003	pyspark	idle	Link	Link	✓

SparkSession available as 'spark'.

Collecting pandas==1.0.3

Using cached https://files.pythonhosted.org/packages/4a/6a/94b219b8ea0f2d580169e85ed1edc0163743f55aaeca8a44c2e8fc1e344e/pandas-1.0.3-cp37-cp37m-manylinux1_x86_64.whl
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas==1.0.3)

Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from pandas==1.0.3)

Collecting python-dateutil>=2.6.1 (from pandas==1.0.3)

Using cached https://files.pythonhosted.org/packages/d4/70/d60450c3dd48ef87586924207ae8907090de0b306af2bce5d134d78615cb/python_dateutil-2.8.1-py2.py3-none-any.whl

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas==1.0.3)

Installing collected packages: python-dateutil, pandas

Successfully installed pandas-1.0.3 python-dateutil-2.8.1

Collecting matplotlib==3.2.1

Using cached https://files.pythonhosted.org/packages/b2/c2/71fcf957710f3balf09088b35776a799ba7dd95f7c2b195ec800933b276b/matplotlib-3.2.1-cp37-cp37m-manylinux1_x86_64.whl

Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1606145839077-0/lib/python3.7/site-packages (from matplotlib==3.2.1)

Collecting pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 (from matplotlib==3.2.1)

Using cached <https://files.pythonhosted.org/packages/8a/bb/488841f56197b13700afd5658fc279a2025a39e22449b7cf29864669b15d/pyparsing-2.4.7-py2.py3-none-any.whl>

Collecting cyclcr>=0.10 (from matplotlib==3.2.1)

Using cached <https://files.pythonhosted.org/packages/f7/d2/e07d3ebb2bd7af696440ce7e754c59dd546ffef1bbe732c8ab68b9c834e61/cyclcr-0.10.0-py2.py3-none-any.whl>

Requirement already satisfied: numpy>=1.11 in /usr/local/lib64/python3.7/site-packages (from matplotlib==3.2.1)

Collecting kiwisolver>=1.0.1 (from matplotlib==3.2.1)

Using cached https://files.pythonhosted.org/packages/d2/46/231de802ade4225b76b96cffe419cf3ce52bbe92e3b092cf12db7d11c207/kiwisolver-1.3.1-cp37-cp37m-manylinux1_x86_64.whl

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib==3.2.1)

Installing collected packages: pyparsing, cyclcr, kiwisolver, matplotlib

Successfully installed cyclcr-0.10.0 kiwisolver-1.3.1 matplotlib-3.2.1 pyparsing-2.4.7

Collecting seaborn==0.10.0

Using cached <https://files.pythonhosted.org/packages/70/bd/5e6bf595fe6ee0f257ae49336dd180768cled3d7c7155b2fdf894c1c808a/seaborn-0.10.0-py3-none-any.whl>

Requirement already satisfied: pandas>=0.22.0 in /mnt/tmp/1606145839077-0/lib/python3.7/site-packages (from seaborn==0.10.0)

Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from seaborn==0.10.0)

Collecting scipy>=1.0.1 (from seaborn==0.10.0)

Using cached https://files.pythonhosted.org/packages/dc/7e/8f6a79b102ca1ea928bae8998b05bf5dc24a90571db13cd119f275ba6252/scipy-1.5.4-cp37-cp37m-manylinux1_x86_64.whl

Requirement already satisfied: matplotlib>=2.1.2 in /mnt/tmp/1606145839077-0/lib/python3.7/site-packages (from seaborn==0.10.0)

Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas>=0.22.0->seaborn==0.10.0)

Requirement already satisfied: python-dateutil>=2.6.1 in /mnt/tmp/1606145839077-0/lib/python3.7/site-packages (from pandas>=0.22.0->seaborn==0.10.0)
 Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /mnt/tmp/1606145839077-0/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn==0.10.0)
 Requirement already satisfied: cyclor>=0.10 in /mnt/tmp/1606145839077-0/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn==0.10.0)
 Requirement already satisfied: kiwisolver>=1.0.1 in /mnt/tmp/1606145839077-0/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn==0.10.0)
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas>=0.22.0->seaborn==0.10.0)
 Installing collected packages: scipy, seaborn
 Successfully installed scipy-1.5.4 seaborn-0.10.0

Importing & Loading Data

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df = spark.read.json('s3://sta9760-project02-dataset/yelp_academic_dataset_business.json')
```

Overview of Data

```
In [3]: print(f'Columns: {len(df.dtypes)} | Rows: {df.count():,}')
```

Columns: 14 | Rows: 209,393

```
In [4]: df.printSchema()
```

```
root
|-- address: string (nullable = true)
|-- attributes: struct (nullable = true)
|   |-- AcceptsInsurance: string (nullable = true)
|   |-- AgesAllowed: string (nullable = true)
|   |-- Alcohol: string (nullable = true)
|   |-- Ambience: string (nullable = true)
|   |-- BYOB: string (nullable = true)
|   |-- BYOBCorkage: string (nullable = true)
|   |-- BestNights: string (nullable = true)
|   |-- BikeParking: string (nullable = true)
|   |-- BusinessAcceptsBitcoin: string (nullable = true)
|   |-- BusinessAcceptsCreditCards: string (nullable = true)
|   |-- BusinessParking: string (nullable = true)
|   |-- ByAppointmentOnly: string (nullable = true)
|   |-- Caters: string (nullable = true)
|   |-- CoatCheck: string (nullable = true)
|   |-- Corkage: string (nullable = true)
|   |-- DietaryRestrictions: string (nullable = true)
|   |-- DogsAllowed: string (nullable = true)
|   |-- DriveThru: string (nullable = true)
|   |-- GoodForDancing: string (nullable = true)
|   |-- GoodForKids: string (nullable = true)
|   |-- GoodForMeal: string (nullable = true)
|   |-- HairSpecializesIn: string (nullable = true)
|   |-- HappyHour: string (nullable = true)
|   |-- HasTV: string (nullable = true)
|   |-- Music: string (nullable = true)
|   |-- NoiseLevel: string (nullable = true)
|   |-- Open24Hours: string (nullable = true)
|   |-- OutdoorSeating: string (nullable = true)
|   |-- RestaurantsAttire: string (nullable = true)
|   |-- RestaurantsCounterService: string (nullable = true)
|   |-- RestaurantsDelivery: string (nullable = true)
|   |-- RestaurantsGoodForGroups: string (nullable = true)
|   |-- RestaurantsPriceRange2: string (nullable = true)
|   |-- RestaurantsReservations: string (nullable = true)
|   |-- RestaurantsTableService: string (nullable = true)
|   |-- RestaurantsTakeOut: string (nullable = true)
|   |-- Smoking: string (nullable = true)
|   |-- WheelchairAccessible: string (nullable = true)
|   |-- WiFi: string (nullable = true)
|-- business_id: string (nullable = true)
|-- categories: string (nullable = true)
```

```

|-- city: string (nullable = true)
|-- hours: struct (nullable = true)
|   |-- Friday: string (nullable = true)
|   |-- Monday: string (nullable = true)
|   |-- Saturday: string (nullable = true)
|   |-- Sunday: string (nullable = true)
|   |-- Thursday: string (nullable = true)
|   |-- Tuesday: string (nullable = true)
|   |-- Wednesday: string (nullable = true)
|-- is_open: long (nullable = true)
|-- latitude: double (nullable = true)
|-- longitude: double (nullable = true)
|-- name: string (nullable = true)
|-- postal_code: string (nullable = true)
|-- review_count: long (nullable = true)
|-- stars: double (nullable = true)
|-- state: string (nullable = true)

```

Display the first 5 rows with the following columns:

- business_id
- name
- city
- state
- categories

```
In [5]: df.select(df['business_id'], df['name'], df['city'], df['state'], df['stars'], df['categories'])
```

business_id	name	city	state	stars	categories
f9NumwFMBDn751xgF...	The Range At Lake...	Cornelius	NC	3.5	Active Life, Gun/...
YzvJg0SayhoZgCljU...	Carlos Santo, NMD	Scottsdale	AZ	5.0	Health & Medical,...
XNoUzKckATkOD1hP6...	Felinus	Montreal	QC	5.0	Pets, Pet Service...
6OAZjbxqM5o129BuH...	Nevada House of Hose	North Las Vegas	NV	2.5	Hardware Stores, ...
51M2Kk903DFYI6gnB...	USE MY GUY SERVIC...	Mesa	AZ	4.5	Home Services, Pl...

only showing top 5 rows

Part 2

Analyzing Categories

- Association Table
- Total Unique Categories
- Top Categories By Business
- Bar Chart of Top Categories

Association Table

```
In [6]: df.createOrReplaceTempView('business')

sqldf = spark.sql(
'''
SELECT *, EXPLODE(SPLIT(categories, ',')) AS category
FROM business
'''
)

sqldf.createOrReplaceTempView('category')

association = spark.sql(
'''
SELECT business_id, category
FROM category
LIMIT 5
'''
)

association.show()
```

business_id	category
f9NumwFMBDn751xgF...	Active Life
f9NumwFMBDn751xgF...	Gun/Rifle Ranges
f9NumwFMBDn751xgF...	Guns & Ammo
f9NumwFMBDn751xgF...	Shopping
YzvJg0SayhoZgCljU...	Health & Medical

Total Unique Categories

```
In [7]: uni_categories_count = spark.sql(
'''
SELECT COUNT(DISTINCT category) as unique_category
FROM category
'''
)

uni_categories_count.show()
```

unique_category
1336

Top Categories By Business

```
In [8]: top_20 = spark.sql(
'''
SELECT category, COUNT(*) AS count
FROM category
GROUP BY category
ORDER BY count DESC
LIMIT 20
'''
)

top_20.show()
```

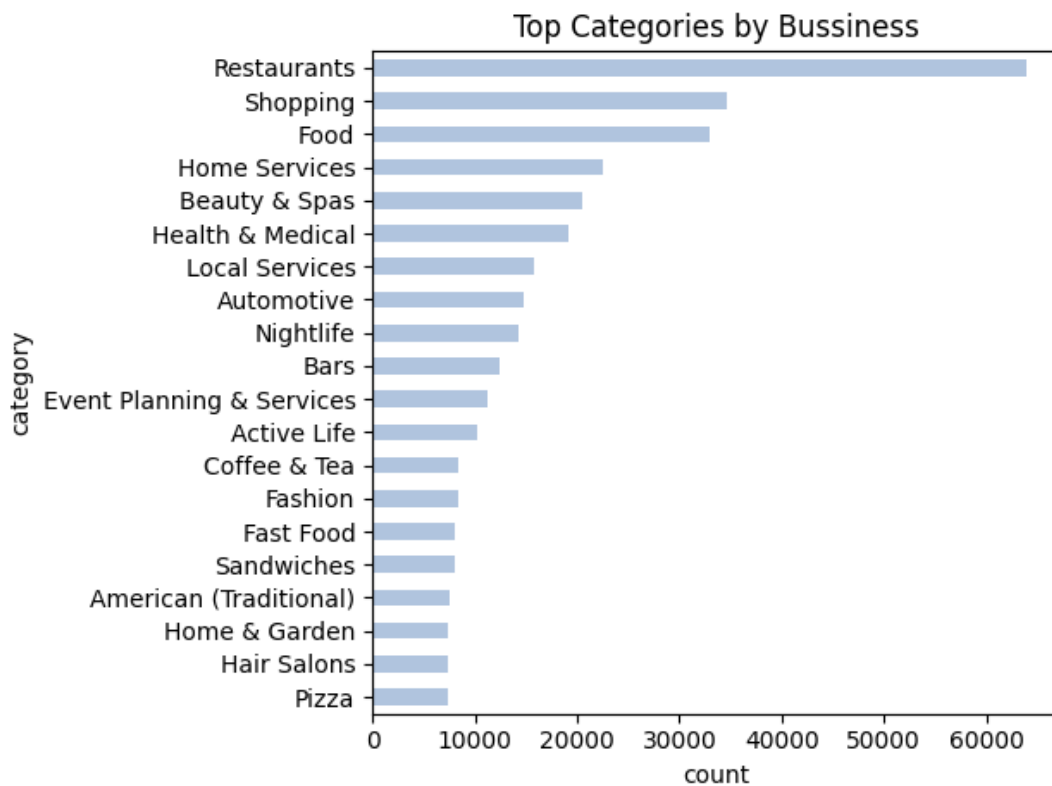
category	count
Restaurants	63944
Shopping	34644
Food	32991
Home Services	22487
Beauty & Spas	20520
Health & Medical	19227
Local Services	15783
Automotive	14720
Nightlife	14211
Bars	12400
Event Planning & ...	11263
Active Life	10225
Coffee & Tea	8415
Fashion	8374
Fast Food	8106
Sandwiches	8064
American (Traditi...	7596
Home & Garden	7331
Hair Salons	7303
Pizza	7302

Bar Chart of Top Categories

```
In [9]: top_20.toPandas().plot.barh(y = 'count', x = 'category', color = 'lightsteelblue', legend=False)
plt.title('Top Categories by Business')
plt.ylabel('category')
plt.xlabel('count')
plt.gca().invert_yaxis()
```

```
plt.tight_layout()

%matplotlib plt
```



Part 3

Do Yelp Reviews Skew Negative?

For this next part, you will attempt to answer the question: are the (written) reviews generally more pessimistic or more optimistic as compared to the overall business rating.

Loading User Data

```
In [10]: df_review = spark.read.json('s3://sta9760-project02-dataset/yelp_academic_dataset_review.json')
df_review.printSchema()
```

```
root
|-- business_id: string (nullable = true)
|-- cool: long (nullable = true)
|-- date: string (nullable = true)
|-- funny: long (nullable = true)
|-- review_id: string (nullable = true)
|-- stars: double (nullable = true)
|-- text: string (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
```

```
In [11]: df_review[df_review['business_id'], df_review['stars']].show(5)
```

```
+-----+-----+
| business_id | stars |
+-----+-----+
|-MhfebM0QIsKt87iD...| 2.0 |
|lbrU8StCq3yDfr-QM...| 1.0 |
|Hq128KMwrEKHqhFrr...| 5.0 |
|5JxlZaqCnk1MnbgRi...| 1.0 |
|IS4cv902ykd8wj1TR...| 4.0 |
+-----+-----+
only showing top 5 rows
```

Rating by users who took the time to submit a written review

```
In [12]: df_review.createOrReplaceTempView('review')

do_review_rating = spark.sql(
'''
SELECT business_id, AVG(stars) AS `avg(stars)`
FROM review
WHERE text IS NOT NULL
GROUP BY business_id
'''
)

do_review_rating.show(5)
```

```
+-----+-----+
|      business_id|      avg(stars)|
+-----+-----+
|RMjCnixEY5i12Ciqn...|3.5316455696202533|
|VHsNB3pdGVcRgs6C3...| 3.411764705882353|
|kpbhERZojleTDRnMV...| 2.0333333333333333|
|ipFreSFhjClfNETuM...|                2.6|
|9A_mB7Ez3RIh26EN5...|                2.6|
+-----+-----+
only showing top 5 rows
```

Join two dataframes (reviews and business data)

```
In [13]: do_review_rating.createOrReplaceTempView('dReview')

review_business_data = spark.sql(
'''
SELECT *
FROM business b
JOIN dReview d
ON b.business_id = d.business_id
'''
)

review_business_data.createOrReplaceTempView("review_business_data")

join_rbd_table = spark.sql(
'''
SELECT `avg(stars)`, stars, name, city, state
FROM review_business_data
ORDER BY `avg(stars)` DESC
'''
)

join_rbd_table.show(5)
```

```
+-----+-----+-----+-----+-----+
|avg(stars)|stars|      name|      city|state|
+-----+-----+-----+-----+-----+
|        5.0| 5.0| Larry Fafalak, LMT| Las Vegas| NV|
|        5.0| 5.0|  Rentech Solutions|Willoughby| OH|
|        5.0| 5.0| Everest Curry and...|  Calgary| AB|
|        5.0| 5.0| Krown Rust Contro...|  Markham| ON|
|        5.0| 5.0|      EVO Swim School|    Mesa|  AZ|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

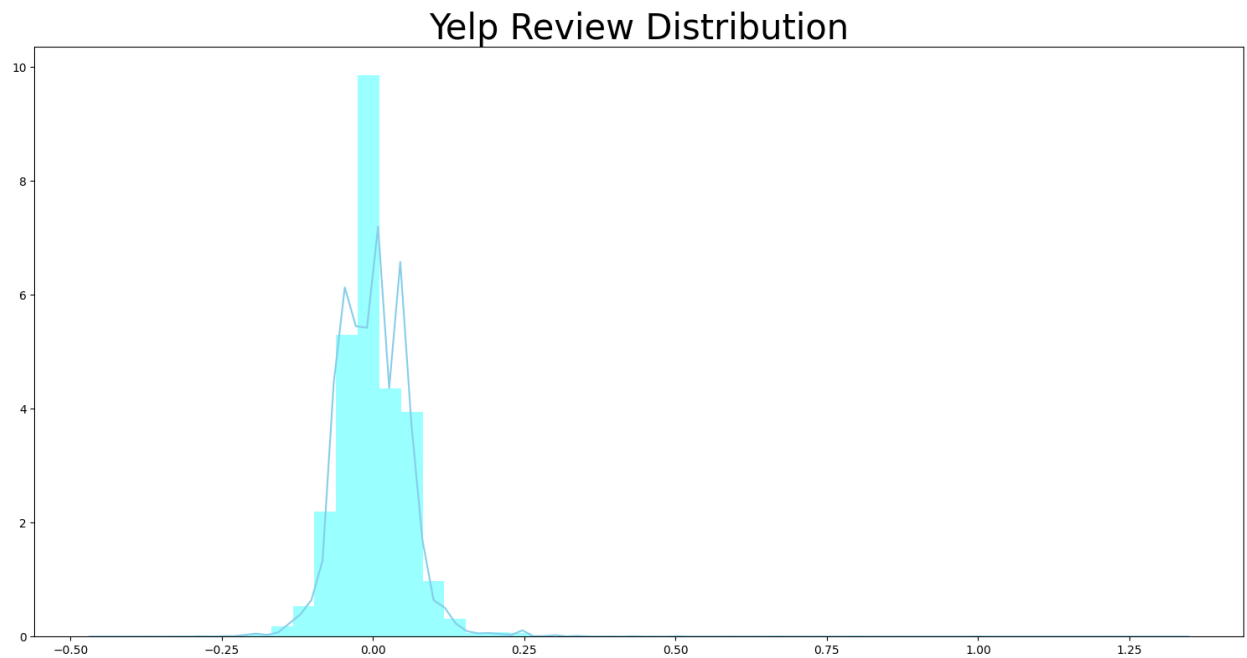
Calculates the skewness

```
In [14]: join_rbd_table.createOrReplaceTempView('join_table')

distri = spark.sql(
'''
SELECT (`avg(stars)` - stars)/ stars AS skew
FROM join_table
'''
)
```

```
In [15]: skew = distri.toPandas()
fig = plt.figure(figsize=(15,8))
plot = sns.distplot(skew, color='skyblue', hist_kws={'color':'cyan'})
plt.title('Yelp Review Distribution', fontsize=30)
plt.tight_layout()

%matplotlib plt
```



```
In [16]: print(f'Total number of review rating data:{df_review.count()}',
            f'Total number of do review rating data:{do_review_rating.count()}', sep='\n')
```

```
Total number of review rating data:8021122
Total number of do review rating data:209393
```

Analysis for part 3

- From the graph above, the different between the average rating and the rating with written review seems to follow a normal distribution, which might indicate that there is no skewness. Besides, the proportion of do review rating data is only about 2.5% of the total number of data, we can ignore the effects that using the average star of all data in the deductions of fraction to calculate the skewness instead of the average star of data with no review. Therefore it's fair to say that the argument that "only people who write a written review are those who are extremely dissatisfied or extremely satisfied with the service received" is not quite true.

Further Discussion

- On the other hand, we know that some reviewers might only leave a very short comment, such as "good", "excellent"...etc. So taking these rating of short reviews as do review rating may have bias.
- Instead of using do review/ not do review rating, I divide the short review and long review by the length of text(50) to make further examination.

```
In [17]: short_review_rating = spark.sql(
    '''
    SELECT business_id, AVG(stars) AS `avg(stars) with short review`
    FROM review
    WHERE LENGTH(text) < 50
    GROUP BY business_id
    '''
)
```

```
long_review_rating = spark.sql(
'''
SELECT business_id, AVG(stars) AS `avg(stars) with long review`
FROM review
WHERE LENGTH(text) > 50
GROUP BY business_id
'''
)
```

```
In [18]: print(f'Total number of short review rating data:{short_review_rating.count()}',
            f'Total number of long review rating data:{long_review_rating.count()}', sep='\n')
```

```
Total number of short review rating data:7740
Total number of long review rating data:209393
```

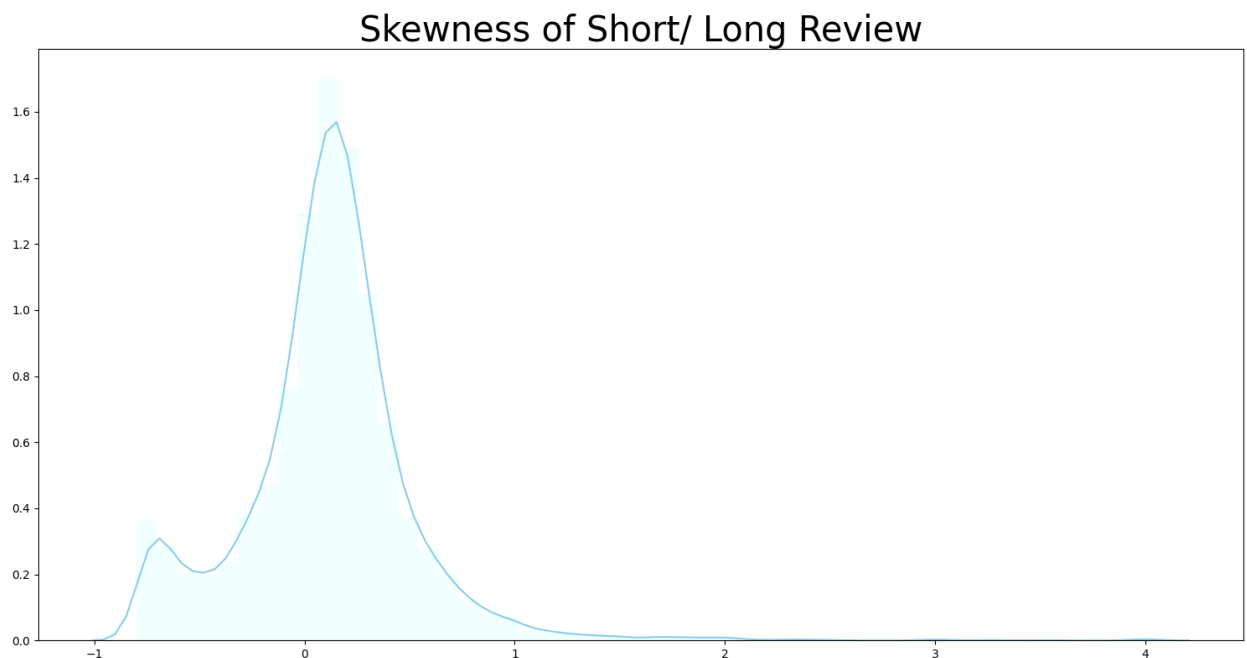
- The proportion of short review rating data of the total number of data is very small, we can ignore the effects.
- I then examine the skew of short/long review rating.

```
In [19]: short_review_rating.createOrReplaceTempView("short_review_rating")
long_review_rating.createOrReplaceTempView("long_review_rating")

skew_2 = spark.sql(
'''
SELECT (`avg(stars) with short review` - `avg(stars) with long review`)/`avg(stars) with lon
FROM short_review_rating s
JOIN long_review_rating l
ON s.business_id = l.business_id
ORDER BY 1 DESC
'''
)
```

```
In [20]: skew_2 = skew_2.toPandas()
fig = plt.figure(figsize=(15,8))
plot = sns.distplot(skew_2, color='skyblue', hist_kws={'color':'lightcyan'})
plt.title('Skewness of Short/ Long Review', fontsize=30)
plt.tight_layout()

%matplotlib plt
```



- The graph above shows that the skew is positive, we can interpret that to be: reviewers who left a short written response gave much higher rating than reviewers who left a long review on average.

Part 4

Should the Elite be Trusted?

Loading User Data

```
In [21]: df_user = spark.read.json('s3://sta9760-project02-dataset/yelp_academic_dataset_user.json')
```

```
In [22]: df_user.printSchema()
```

```
root
|-- average_stars: double (nullable = true)
|-- compliment_cool: long (nullable = true)
|-- compliment_cute: long (nullable = true)
|-- compliment_funny: long (nullable = true)
|-- compliment_hot: long (nullable = true)
|-- compliment_list: long (nullable = true)
|-- compliment_more: long (nullable = true)
|-- compliment_note: long (nullable = true)
|-- compliment_photos: long (nullable = true)
|-- compliment_plain: long (nullable = true)
|-- compliment_profile: long (nullable = true)
|-- compliment_writer: long (nullable = true)
|-- cool: long (nullable = true)
|-- elite: string (nullable = true)
|-- fans: long (nullable = true)
|-- friends: string (nullable = true)
|-- funny: long (nullable = true)
|-- name: string (nullable = true)
|-- review_count: long (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
|-- yelping_since: string (nullable = true)
```

```
In [23]: df_user.select(df_user['user_id'],df_user['elite']).show(5)
```

```
+-----+-----+
|          user_id|          elite|
+-----+-----+
|ntlvfPzc8eglqvK92...|
|FOBRPlBHa3WPHFB5q...|2008,2009,2010,20...|
|zZUnPeh2hEp0WydbA...|2010|
|QaELAmRcDc5TfJEyl...|2009|
|xvu8G900tezTzbbfq...|2009,2010,2011,20...|
+-----+-----+
only showing top 5 rows
```

Number of Elite

```
In [24]: df_user.createOrReplaceTempView("user")
```

```
num_E = spark.sql(
    '''
    SELECT COUNT(*) AS `number of elite`
    FROM user u
    JOIN review r
    ON u.user_id = r.user_id
    WHERE elite LIKE '%20%'
    '''
)

num_E.show()
```

```
+-----+
|number of elite|
```

```
+-----+
|      1756327      |
+-----+
```

Avg. Rating from Elite

```
In [25]: avgR_from_E = spark.sql(
'''
SELECT SUM(average_stars)/ COUNT(*) AS `avgstars from elite`
FROM user u
JOIN review r
ON u.user_id = r.user_id
WHERE elite LIKE '%20%'
'''
)

avgR_from_E.show()
```

```
+-----+
|avgstars from elite|
+-----+
| 3.8541299199978103|
+-----+
```

Avg. Rating from Pedestrian

```
In [26]: avgR_from_P = spark.sql(
'''
SELECT SUM(average_stars)/ COUNT(*) AS `avgstars from pedestrian`
FROM user u
JOIN review r
ON u.user_id = r.user_id
WHERE elite NOT LIKE '%20%'
'''
)

avgR_from_P.show()
```

```
+-----+
|avgstars from pedestrian|
+-----+
| 3.68831514838085|
+-----+
```

```
In [27]: # Join tables
df_elite_user = spark.sql(
'''
SELECT r.business_id, AVG(r.stars) AS `avg.star from Elite`
FROM user u
JOIN review r
ON u.user_id = r.user_id
WHERE elite LIKE '20%'
GROUP BY r.business_id
'''
)

df_ped_user = spark.sql(
'''
SELECT r.business_id, AVG(r.stars) AS `avg.star from Pedestrian`
FROM user u
JOIN review r
ON u.user_id = r.user_id
WHERE elite NOT LIKE '20%'
GROUP BY r.business_id
'''
)
```

Rating difference

```
In [28]: df_elite_user.createOrReplaceTempView("elite_user")
```

```
df_ped_user.createOrReplaceTempView("ped_user")

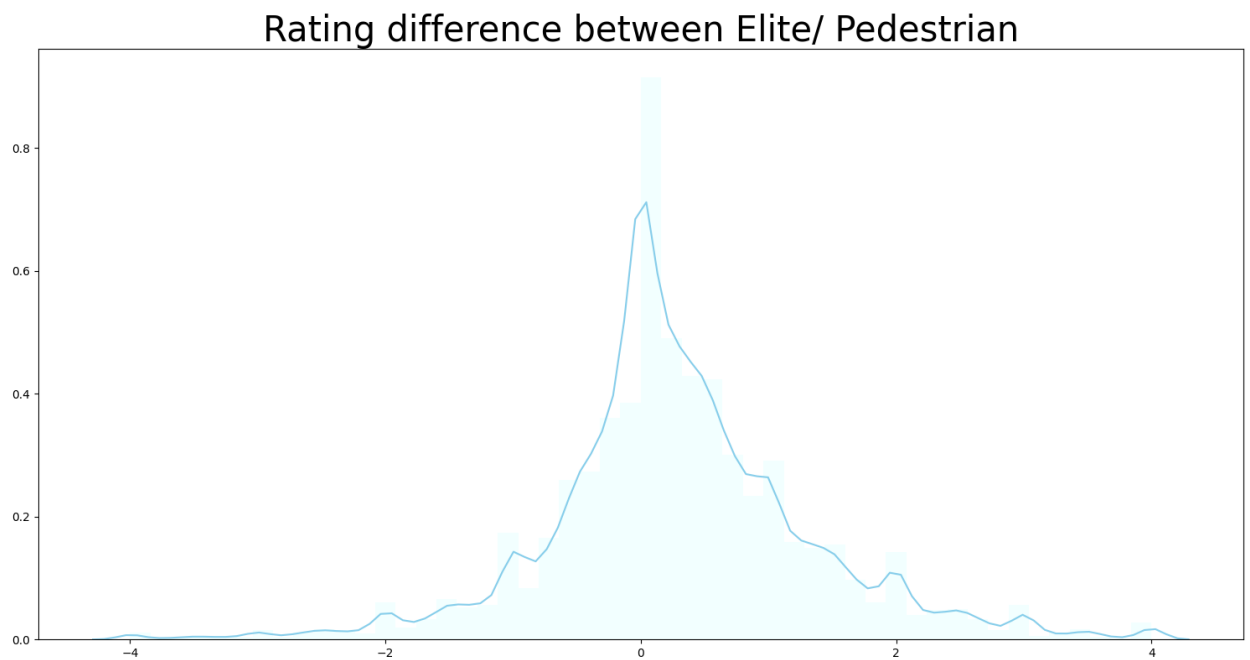
diff = spark.sql(
    '''
    SELECT (`avg.star from Elite` - `avg.star from Pedestrian`) AS diff
    FROM elite_user e
    JOIN ped_user p
    ON e.business_id = p.business_id
    ORDER BY 1 DESC
    '''
)

diff.show(5)
```

```
+----+
|diff|
+----+
| 4.0|
| 4.0|
| 4.0|
| 4.0|
| 4.0|
+----+
only showing top 5 rows
```

```
In [29]: diff = diff.toPandas()
fig = plt.figure(figsize=(15,8))
plot= sns.distplot(diff, color='skyblue',hist_kws={'color':'lightcyan'})
plt.title('Rating difference between Elite/ Pedestrian', fontsize=30)
plt.tight_layout()

%matplotlib plt
```



```
In [30]: print(diff.describe())
```

```

              diff
count  146479.000000
mean      0.296621
std       1.044539
min       -4.000000
25%      -0.214286
50%       0.195538
75%       0.811914
max        4.000000
```

- From the graph above, there are not much different between the average rating from elite and pedestrian.
- The avg.star from elite is slightly higher than from pedestrian. ###

- According to these two results above, I think we can trust elite user although sometimes they might give a little higher rating than pedestrian. The possible reason for this slightly deviation is that some elite might have cooperation with the business, the rating is just a result of advertorial. Nonetheless, the deviation is acceptable for me. If someone really care about this deviation, they can lower the rating a bit from elite on themself.

Other Analysis

Effects on fans number from number of review/ useful review

```
In [31]: effects_on_fan_num_review = spark.sql(
'''
SELECT SUM(fans)/ SUM(review_count) AS `effect on fan from number of review`
FROM user
'''
)

effects_on_fan_num_review.show()
```

```
+-----+
|effect on fan from number of review|
+-----+
|                0.06580379113050557|
+-----+
```

```
In [32]: effects_on_fan_useful_review = spark.sql(
'''
SELECT SUM(fans)/ SUM(useful) AS `effect on fan from useful of review`
FROM user
'''
)

effects_on_fan_useful_review.show()
```

```
+-----+
|effect on fan from useful of review|
+-----+
|                0.03662893096969366|
+-----+
```

- From the results above, it seems that the user should focus on increasing the number of the review rather than focus on the quality of review if they want to efficiently increase the fans number.

In []: